Where the grammar's advantage actually is: not raw search power, but validity and compactness

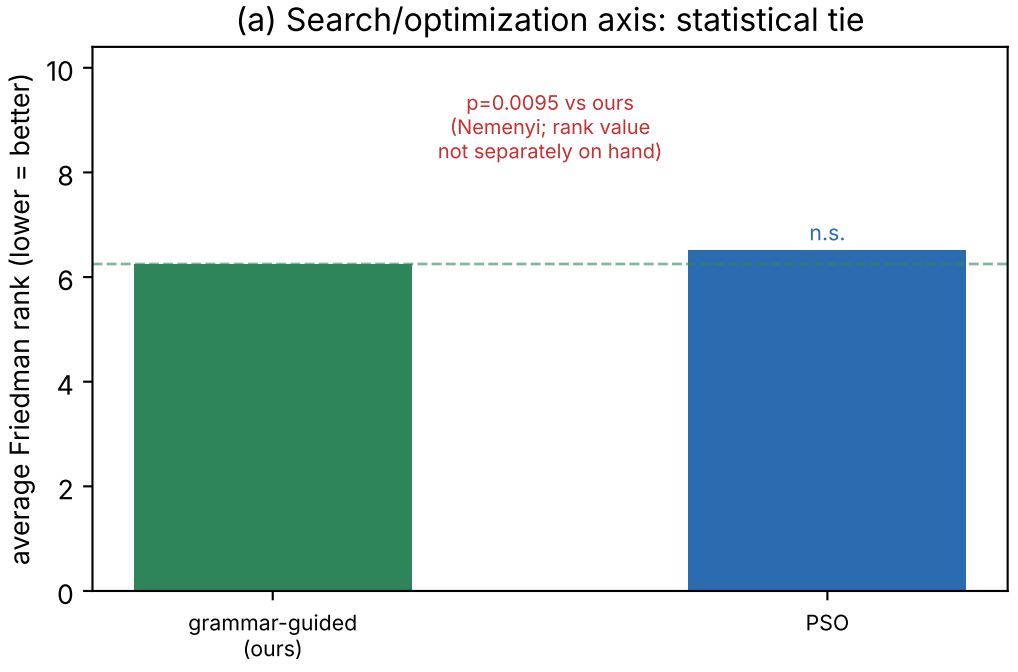

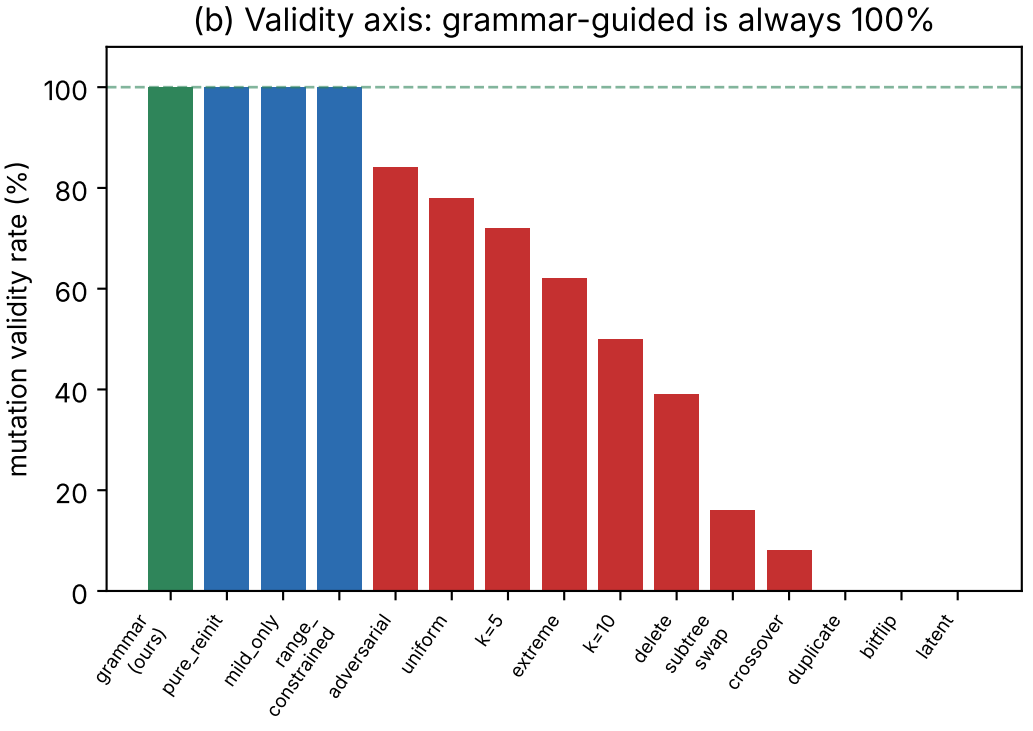

Only grammar-guided (6.25, best of 21) and PSO (6.50) have a confirmed average-rank value on hand from prior output. The other 18 methods, including the full CMA-ES family, are tied with ours per Nemenyi (p>0.05) but their exact rank values aren't re-plotted here without re-reading expanded_baseline_table.py's saved output. CMA-ES lite vs. standard: separate paired t-test p=0.816.

Compactness (reviewer3_compactness_curve.py): 6 typed rules at every body size vs. 17/37/49 flat rules (4×2/6×3/8×3 legs x segments) -- 2.8x/6.2x/8.2x, up to 13.5x in the full sweep.